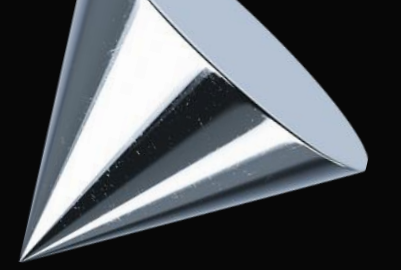


# Assignments

**Navdeep Kaur**

# Dockers Lab1



# Duration - 20 mins

1. Run a container with below parameters

## Image Info

Image: `navjoy220161/python_flask:1.0.0`

## Container Info

Container name: `python-c`

Mode: `Detached`

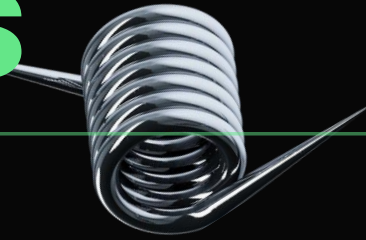
Forward 8087 of host -> port 5000 of container

2. Check the output on host:8087
3. Now enter inside the container
  - `rm` the file `main.py`
4. Now stop and start the container
5. Access the page `localhost:8087`
6. Check container logs and status
7. Remove the container
8. Create the container again using the same command as in step1
9. Change the content of your local `main.py` file to display "Hello Students"
10. Replace the container's `main.py` with your local edited `main.py` and check the page again



# Dockers Lab2



# Duration - 20 mins

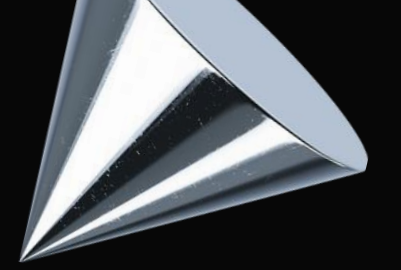


1. Create a folder mypython-image and copy paste the main.py from python-flask-demo
2. Rename main.py to entry.py
3. Edit the content of entry.py to display "This is my python image"
4. Create a docker file with below steps
  - from** image navjoy220161/python\_flask:1.0.0
  - remove** main.py from /python-flask
  - copy** entry.py created by you inside the container
  - Install** pandas libraryFinally run the command "python entry.py" as an entrypoint of the container
5. Create a DockerHub Account
6. Run docker login
7. Create an image from above docker file with tag mypython:1.0.0
8. Check your repository in dockerhub
9. Remove the image locally
10. Run a container from above image and expose the port 5000 on 8089 on host
11. Check localhost:8089




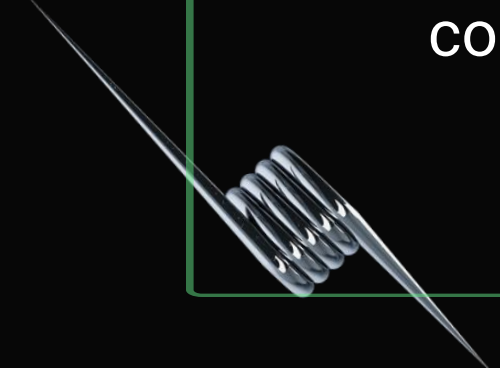


# Pods Assignment





# Duration - 20 mins

1. Create a new pod with the nginx image and name nginx-pod without using a yaml file.
  2. Create a new pod using yaml file with below configuration :
    - pod name : nginx-pod-1
    - Label: type=loadbalancer, country=us
    - Container name: nginx-container
    - Container image: nginx, tag 1.21.1-alpine
  3. There is a Pod named httpd-pod that is already in your machine. There is some issue while running that pod. Fix the issue and make it in Running state. Also change the Pod label to frontend  
(Hint: get yaml file from running pod, delete existing pod, change yaml file and create new pod using `kubectl create -f <name>.yaml`)
  4. There is a pod named myweb-app that is already running in your machine. That pod contains two containers - httpd-container and mongodb-container.
    - >> Get the logs of container mongodb-container
    - >> Go inside httpd-container and execute the command `whoami`
- 
- 

# Replica Set Assignment





# Duration - 20 mins

1. Create below replicaset using yaml file with below configuration :

Replicaset name: nginx-rs

Labels: type:loadbalancer-replica

Container spec

pod name : nginx-pod

Label: type=loadbalancer, country=us

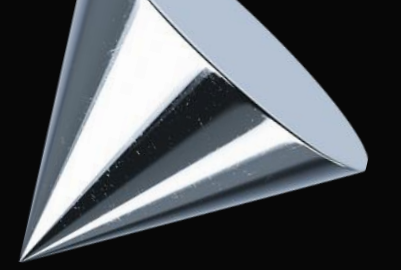
Container name: nginx-container

Container image: nginx, tag 1.21.1-alpine

Replicas: 3

2. Delete one of the pod created by replicaset and check the pod status
3. There is an existing replicaset named web-rs running in your machine. First check how many replicas it is running right now and scale it up to run 6 replicas
4. There is an existing replicaset named loadbalancer-rs running in your machine. Pod started by this replicas are in error state, fix the issue and make the pod in running state.

# Namespace

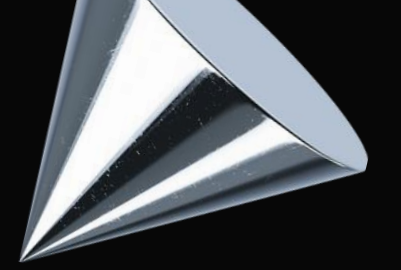


# Duration - 10 mins

1. Check all existing namespace in your machine
2. Checkout the pods that are running in namespace mars
3. Create a new namespace Jupiter and run a pod inside it with an image busybox
4. Check all the pods running across all the namespace


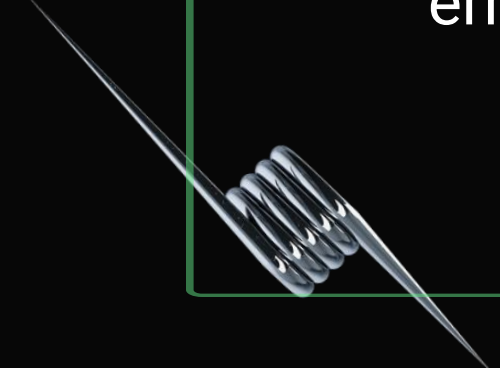


# Deployment Assignment





# Duration - 20 mins

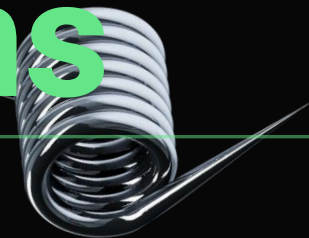
1. Create a new deployment with below configuration :  
deployment name : nginx-deployment  
label: country=us  
Container name: nginx-container  
Container image: nginx, tag 1.21.1-alpine  
Container label: type=loadbalancer  
Replicas: 3
  2. In above deployment , set the image to nginx-junk and check the status
  3. Now check the history of all rollouts
  5. Now rollback to previous version and check status
  6. There is an existing deployment in your machine named myweb-dep whose pods are in error state.
    - . First checkout all the revision of this deployment
    - . Now rollback to a revision where image version was 2.4
- 
- 


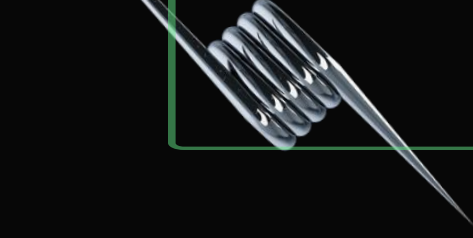
# ConfigMaps



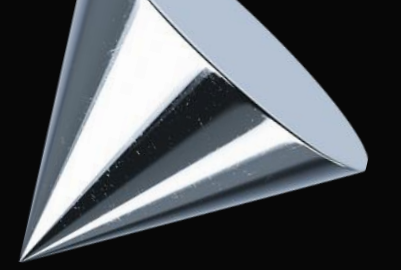


# Duration - 20 mins



1. Create a configmap for a gaming app using command with below properties  
name: gameconfigmap  
data:  
enemies=aliens  
lives=3
  2. Create a second configmap for the same app but this time using yaml.  
name: uiconfigmap  
data:  
theme=black  
badplayer=red  
goodplayer=purple
  3. Create a pod definition and inject the configs from above two configMap in it.  
pod name – httpd-game-pod  
image – httpd  
ENEMIES\_TYPE -> map config from gameconfigmap with key “enemies”  
TOTAL\_LIVES -> map config from gameconfigmap with key “lives”  
Also Inject the whole uiconfigmap at one go
- 
- 

**Secrets**



# Duration - 25 mins

1. Create a secret for the same gaming app using command with below properties

name: dbsecret

data:

db\_host=mongoDB

db\_user=root

db\_pass=mongo123

3. Now Create a second secret for the same gaming app using yaml file

name: gamesecret

data:

game\_user= admin

game\_pass=admin123

3. Now create a pod with name httpd-pod and image httpd and inject secret

MONGO\_HOST -> db\_host from secret dbsecret

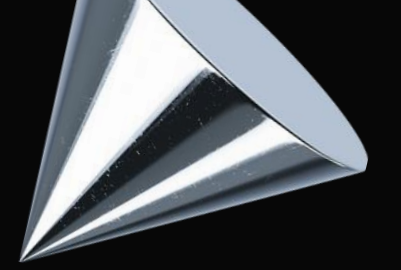
MONGO\_USER -> db\_user from secret dbsecret

MONGO\_PASS -> db\_pass from secret dbsecret

Inject gamesecret as it is

4. Modify the existing pod named httpd-pod and mount the gamesecret as a volume mounted at /opt/secret

# Node-Mongo Project



# Duration - 15 mins



Before doing this assignment, make sure your cluster is clean

1. Go to below location

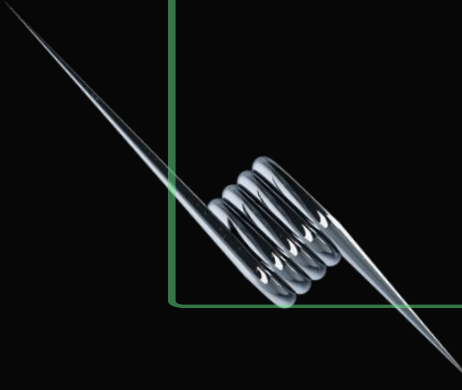

`/Kubernetes_Fundamentals/k8s_operations/deployment/node-mongo`

2. Run below command

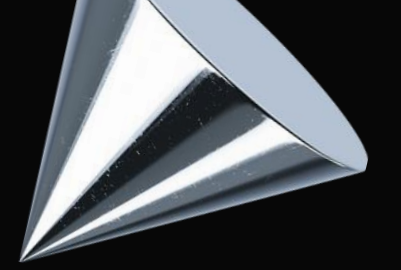
`kubectl create -f .`

`Kubectl rollout restart deployment mywebapp`

3. Now explore your cluster for below answers
- 

1. How many deployment objects are there?
  2. Checkout webapp deployment- how many replicas, what is the image used?
  3. Which Service is exposing webapp deployment - what is the type, name and matching label of that service?
  4. Checkout mongodb deployment- how many replicas, what is the image used?
  5. Which Service is exposing mongodb deployment - what is the type, name and matching label of that service?
  6. Also check out configmap that got created and what is Mongo\_host referred too here?
  7. In which deployment, configmap is injected?
  8. Is APP\_port of configmap and webapp-service matching?
- 
- 

# Persistent Volumes







# Duration - 15 mins

1. Check out what is default storage class in your cluster.

```
kubectl get storageclass
```

2. Describe storage class "gp2"

```
kubectl describe storageclass gp2
```

3. Create a persistent volume claim using the storage class "gp2" and claim 10Gi. You can change below yaml file

```
/Kubernetes_Fundamentals/k8s_operations/deployment/node-mongo-persistence/mongo-pvc.yaml
```

4. Check your pvc

```
kubectl get pvc
```

5. Now create mongodb-deployment using above persistent volume claim

```
/Kubernetes_Fundamentals/k8s_operations/deployment/node-mongo-persistence/mongodb-deployment.yaml
```

6. Describe your mongo pod check whether it is using above pvc?

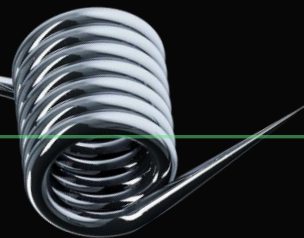
```
kubectl describe pod mongo-58c79fcb66-kqmwv
```



**DaemonSet**



# Duration - 15 mins



1. Go to below location

[/Kubernetes\\_Fundamentals/k8s\\_operations/DaemonSet](#)

2. Check out the daemonset.yaml

3. Did you notice there is no replicas in the spec part

4. Create the Daemonset

`Kubectl create -f daemonset.yaml`

5. Check out the pod

`Kubectl get pod`



7. Did you noticed, you got exactly two replicas because if you check you have two worker machine

8 . Check out NODE in result of above command

9 . Check out node of your cluster

`Kubectl get node`

10 . Check out pod and node on which they are allocated

`Kubectl get pod`

11. Did you noticed, you got exactly two replicas on two different nodes

# Canary Deployment



# Duration - 15 mins

1. Create your first deployment with below configuration :

deployment name : webserver-deployment

label:

app: frontend

version: v1

image: httpd

Replicas: 3

2. Above deployment is exposed on port 80. Add a NodePort service on nodePort 30000 to access above deployment

3. Let's hit the service multiple times

```
for ((i=1;i<=10;i++)); do curl 192.168.49.2:30000; done
```

4. Now you want to deploy a new version with different image. You want to deploy it using canary approach.

Create a second deployment such that 25% of traffic comes to the new version and the rest 75% goes to older version v1

Image : nginx

Replicas: 1

5. Let's hit the service multiple times again

```
for ((i=1;i<=10;i++)); do curl 192.168.49.2:30000; done
```

Hint: Use labels in such a way that existing service starts sending traffic to new deployment as well

# **Blue-Green Deployment**





# Duration - 15 mins



1. Create your first deployment with below configuration :

deployment name : webserver-deployment-blue

label:

version: blue

image: httpd

Replicas: 3

2. Above deployment is exposed on port 80. Add a NodePort service on nodePort 30000 to access above deployment

3. Let's hit the service multiple times

for ((i=1;i<=10;i++)); do curl 192.168.49.2:30000; done

4. Now you want to deploy a new version. You want to deploy it using blue/green approach.

Create a second deployment with same number of replicas as your existing deployment

deployment name : webserver-deployment-green

label:

version: green

image: nginx

Replicas: 3

5. Adjust the labels of your existing service so that it start moving traffic at once to new deployment

6. Let's hit the service multiple times again

for ((i=1;i<=10;i++)); do curl 192.168.49.2:30000; done

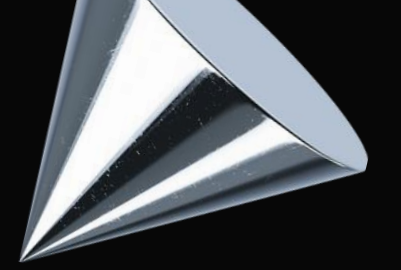
# Taints & Tolerations



# Duration - 15 mins

1. Run a pod on your machine using below command & make sure pod comes in running state  
`kubectl run nginx-pod - --image nginx`
2. Now taint your machine with below taint  
`kubectl taint node minikube env=prod:NoExecute`
3. Now check out the status of pod that you ran in step1  
`kubectl get pod`
4. You might have noticed that pod is in pending status. Check out why by looking at the event section.  
`kubectl describe pod nginx-pod`
5. Now, run a new pod with image "httpd" using yaml file and add matching toleration to run it on tainted machine.
6. Check out the status of above pod. It should be in running status.
7. Let's add a new untainted machine(node) in your minikube cluster so that pod that was missing the toleration should also start running on this new machine  
`minikube node add`
8. Check the status of that nginx-pod that was in pending state. Did it start running, check out on which machine  
`kubectl get pod --o wide`
9. Finally don't forget to untaint your machine  
`kubectl taint node minikube env=prod:NoExecute-`

**Node Affinity / Anti-Affinity**



# Duration – 15 mins

Please note: In the last assignment, you already added one more machine in your minikube cluster. Make sure you have two machines and none of them is tainted.

1. Create a pod with an image mongo. Add a node affinity section to make it run on a machine that has label "size=large"
2. Check out the pod status and describe and look out for the reason , why are they in pending state?
3. You might find out that it must be searching for machine with the matching label. Let's label our only machine

```
kubectl label node minikube size=large
```

4. Check out the pod status and did they start running? Also check did they started on the labelled machine?

```
kubectl get pod -o wide
```

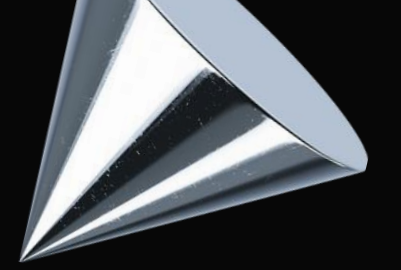
5. What will happen to a pod that does not have node affinity added? Will it run on labelled or unlabelled machine? Try this out, create a deployment with 6 replicas , figure out where those 6 replicas are started?

Did you noticed, replicas are distributed on both machines

6. Modify the above deployment to add anti affinity rules against the minikube master node that is labelled "size=large". Restart the deployment and check out the pod again, where are they started?

Did you noticed, replicas are placed on machine which is not labelled


**Pod Affinity/Anti-Affinity**





# Duration - 15 mins

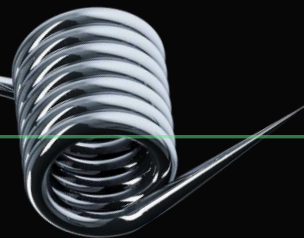
1. We have a frontend app with image "httpd" and cache app with image "redis"
2. frontend pod should not be co-located and neither redis-pods should be co-located.
3. However, Frontend should be deployed alongside redis pod to reduce latency




# Stateful Sets



# Duration - 15 mins



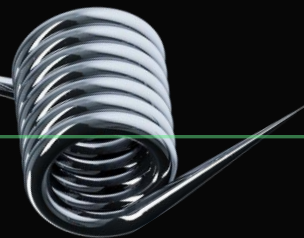
1. Create a StatefulSet for Mysql database with 3 replicas
  2. In your stateful set, use persistence volume claim template to create pvc for 1 GB each using default storage class
  3. Run your Stateful set, check the status of running pods
  4. Check out what all things are created in cluster – Pods, replicaset, services and pvc that is tied to each pod
  - 5 Delete first pod and see whether name and pvc is attached back
  - 6 Try scaling up replicas to 5 , check out the order in which they are started, also check out corresponding pvc
  - 7 Try scaling down replicas to 6 , check out the order in which they are scaled down and also check corresponding pvc
- 

# StatefulSets

Using Kubernetes Operator/ Helm



# Duration - 15 mins



Create a StatefulSet for prometheus/Grafana stack using Kubernetes Operator/Helm

